



Session 19: RDD DEEP DIVE

Assignment 19

Contents

Dataset	2
Statement	2
1:	2
2:	2
3:	2
Spark Operations - Problem Statement	
1	3
Task 1.1 Read the text file, and create a tupled rdd.	3
Task1.2 - Find the count of total number of rows present	3
Task1.3 - What is the distinct number of subjects present in the entire school?	3
Task1.4 - What is the count of the number of students in the school, whose name is Mathew and marks is	
55	4
Spark Operations - Problem Statement	
2	6
Task 2.1 - What is the count of students per grade in the school?	6
Task 2.2 - Find the average of each student (Note - Mathew is grade-1, is different from Mathew in some other grade!)	6
Task 2.3 - What is the average score of students in each subject across all grades?	7
Task 2.4 - What is the average score of students in each subject per grade?	9

Task 2.5 - for all students in grade-2, how many have average score greater than	<u>1</u>
50?	<u>1</u>
Spark Operations – Problem Statement	<u>1</u>
3	<u>3</u>



Dataset

Given a dataset of college students as a text file (name, subject, grade, marks):

```
1 Mathew,science,grade-3,45,12
2 Mathew,history,grade-2,55,13
3 Mark,maths,grade-2,23,13
4 Mark,science,grade-1,76,13
5 John,history,grade-1,14,12
6 John,maths,grade-2,74,13
7 Lisa,science,grade-1,24,12
8 Lisa,history,grade-3,86,13
9 Andrew,maths,grade-1,34,13
10 Andrew,science,grade-3,26,14
11 Andrew,history,grade-1,74,12
12 Mathew,science,grade-2,55,12
13 Mathew,history,grade-2,87,12
14 Mark,maths,grade-1,92,13
15 Mark,science,grade-2,12,12
16 John,history,grade-1,67,13
17 John,maths,grade-1,35,11
18 Lisa,science,grade-2,24,13
19 Lisa,history,grade-2,98,15
20 Andrew,maths,grade-1,23,16
21 Andrew,science,grade-3,44,14
22 Andrew,history,grade-2,77,11
```

Problem Statement

Problem Statement 1:

1. Read the text file, and create a tupled rdd.
2. Find the count of total number of rows present.
3. What is the distinct number of subjects present in the entire school
4. What is the count of the number of students in the school, whose name is Mathew and Marks is 55

Problem Statement 2:

1. What is the count of students per grade in the school?
2. Find the average of each student (Note - Mathew is grade-1, is different from Mathew in some other grade!)
3. What is the average score of students in each subject across all grades?
4. What is the average score of students in each subject per grade?
5. For all students in grade-2, how many have average score greater than 50?

Problem Statement 3:

Are there any students in the college that satisfy the below criteria:

1. Average score per student_name across all grades is same as average score per student_name per grade

Hint - Use Intersection Property.



Spark Operations – Problem Statement 1

Task 1.1 Read the text file, and create a tupled rdd.

Below is the code we use to read the text file using spark context and creating a tuple RDD,

```
scala> val baseRDD =  
sc.textFile("/home/acadgild/hadoop/17.2 Dataset.txt").map(x =>  
(x.split(",")(0),(x.split(",")(1),x.split(",")(2),x.split(",")(3).toInt,x.split(",")  
(4).toInt)))
```

```
scala> baseRDD.foreach(println)
```

We have create a tuple RDD with name as Key and the subject, grades and the marks as values.

```
scala> baseRDD.foreach(println)  
(Mathew,(science,grade-3,45,12))  
(Mathew,(history,grade-2,55,13))  
(Mark,(maths,grade-2,23,13))  
(Mark,(science,grade-1,76,13))  
(John,(history,grade-1,14,12))  
(John,(maths,grade-2,74,13))  
(Lisa,(science,grade-1,24,12))  
(Lisa,(history,grade-3,86,13))  
(Andrew,(maths,grade-1,34,13))  
(Andrew,(science,grade-3,26,14))  
(Andrew,(history,grade-1,74,12))  
(Mathew,(science,grade-2,55,12))  
(Mathew,(history,grade-2,87,12))  
(Mark,(maths,grade-1,92,13))  
(Mark,(science,grade-2,12,12))  
(John,(history,grade-1,67,13))  
(John,(maths,grade-1,35,11))  
(Lisa,(science,grade-2,24,13))  
(Lisa,(history,grade-2,98,15))  
(Andrew,(maths,grade-1,23,16))  
(Andrew,(science,grade-3,44,14))  
(Andrew,(history,grade-2,77,11))
```

Task1.2 - Find the count of total number of rows present

```
scala> baseRDD.count()
```

```
res1: Long = 22
```

By using count() function we can see the number of lines present in the text file.

```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2 Dataset.txt").map(x => (x.split(",")(0),(x.split(",")(1),x.split(",")(2),x.split(",")  
(3).toInt,x.split(",")(4).toInt)))  
baseRDD: org.apache.spark.rdd.RDD[(String, (String, String, Int, Int))] = MapPartitionsRDD[5] at map at <console>:24  
scala> baseRDD.count()  
res1: Long = 22
```

Task1.3 - What is the distinct number of subjects present in the entire school?

Please see the below codes used for this task,

```
val baseRDD =  
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>  
(x.split(",")(1),1))  
  
val RDDreduce = baseRDD.reduceByKey((x,y)=>(x+y))
```



ACADGILD

scala> RDDreduce.foreach(println)**(maths,6)****(history,8)****(science,8)**

First we are creating a RDD to read the file and selecting only subject name and mapping them with value 1 and counting the values of occurrences using **reduceByKey** to get distinct number of subjects.

```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=> (x.split(",")(1),1))
baseRDD: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[169] at map at <console>:24

scala> baseRDD.foreach(println)
(science,1)
(history,1)
(maths,1)
(science,1)
(history,1)
(maths,1)
(science,1)
(history,1)
(maths,1)
(science,1)
(history,1)
(science,1)
(history,1)
(maths,1)
(science,1)
(history,1)
(maths,1)
(science,1)
(history,1)
(maths,1)
(science,1)
(history,1)

scala> val RDDreduce = baseRDD.reduceByKey((x,y)=>{x+y})
RDDreduce: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[170] at reduceByKey at <console>:26

scala> RDDreduce.foreach(println)
(maths,6)
(history,8)
(science,8)
```

Task1.4 - What is the count of the number of students in the school, whose name is Mathew and marks is 55

Below are the codes,

val baseRDD =**sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x => ((x.split(",")(0),x.split(",")(3).toInt),1))****val RDDfilter = baseRDD.filter(x=>x._1._1 == "Mathew" && x._1._2 == 55)****val RDDreduce = RDDfilter.reduceByKey((x,y)=> x+y).foreach(println)**

In the first line code, we are reading the text file and creating a tuple RDD as "baseRDD" with name & marks as key and mapping numerical 1 as value.



```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x => ((x.split(",")(0),x.split(",")(3).toInt),1))
baseRDD: org.apache.spark.rdd.RDD[(String, Int), Int] = MapPartitionsRDD[176] at map at <console>:24

scala> baseRDD.foreach(println)
((Mathew,45),1)
((Mathew,55),1)
((Mark,23),1)
((Mark,76),1)
((John,14),1)
((John,74),1)
((Lisa,24),1)
((Lisa,86),1)
((Andrew,34),1)
((Andrew,26),1)
((Andrew,74),1)
((Mathew,55),1)
((Mathew,87),1)
((Mark,92),1)
((Mark,12),1)
((John,67),1)
((John,35),1)
((Lisa,24),1)
((Lisa,98),1)
((Andrew,23),1)
((Andrew,44),1)
((Andrew,77),1)
```

Filter the tuple RDD by providing the condition Mather and mark as 55,

```
scala> val RDDfilter = baseRDD.filter(x=>x._1._1 == "Mathew" && x._1._2 == 55)
RDDfilter: org.apache.spark.rdd.RDD[(String, Int), Int] = MapPartitionsRDD[177] at filter at <console>:26

scala> RDDfilter.foreach(println)
((Mathew,55),1)
((Mathew,55),1)
```

Now we are counting each occurrences using the **reduceByKey** and the required output is below,

```
scala> val RDDreduce = RDDfilter.reduceByKey((x,y)=> x+y).foreach(println)
((Mathew,55),2)
RDDreduce: Unit = ()
```




Spark Operations – Problem Statement 2

Task 2.1 - What is the count of students per grade in the school?

```
val baseRDD =  
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x =>  
(x.split(",")(2),1)).reduceByKey((x,y)=>x+y).foreach(println)
```

we are reading the text file by creating a tuple RDD with grade as key and mapping numerical 1 as values and reducing the number occurrences using reduceByKey, please see the required output below,

```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x => (x.split(",")(2),1)).reduceByKey((x,y)=>x+y).foreach(println)
(grade-3,4)
(grade-1,9)
(grade-2,9)
baseRDD: Unit = ()
```

Task 2.2 - Find the average of each student (Note - Mathew is grade-1, is different from Mathew in some other grade!)

Codes,

```
val baseRDD =  
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")  
(0),x.split(",")(2)),x.split(",")(3).toInt))
```

```
val RDDmap = baseRDD.mapValues(x=>(x,1))
```

```
val RDDreduce = RDDmap.reduceByKey((x,y) => (x._1 + y._1, x._2 + y._2))
```

```
val StudAvg =  
RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}
```

First we are creating the **baseRDD** to read the file and selecting name and grade as key and marks as value,

```
scala> val RDDmap = baseRDD.mapValues(x=>(x,1))
RDDmap: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = MapPartitionsRDD[186] at mapValues at <console>:26

scala> RDDmap.foreach(println)
((Mathew,grade-3),(45,1))
((Mathew,grade-2),(55,1))
((Mark,grade-2),(23,1))
((Mark,grade-1),(76,1))
((John,grade-1),(14,1))
((John,grade-2),(74,1))
((Lisa,grade-1),(24,1))
((Lisa,grade-3),(86,1))
((Andrew,grade-1),(34,1))
((Andrew,grade-3),(26,1))
((Andrew,grade-1),(74,1))
((Mathew,grade-2),(55,1))
((Mathew,grade-2),(87,1))
((Mark,grade-1),(92,1))
((Mark,grade-2),(12,1))
((John,grade-1),(67,1))
((John,grade-1),(35,1))
((Lisa,grade-2),(24,1))
((Lisa,grade-2),(98,1))
((Andrew,grade-1),(23,1))
((Andrew,grade-3),(44,1))
((Andrew,grade-2),(77,1))
```

We are using reduceByKey to add the occurrences of marks for each key which is student name and grade,



```
scala> val RDDreduce = RDDmap.reduceByKey((x,y) => (x._1 + y._1, x._2 + y._2))
RDDreduce: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = ShuffledRDD[187] at reduceByKey at <console>:28
ACADGILD
scala> RDDreduce.foreach(println)
((Lisa,grade-1),(24,1))
((Mark,grade-2),(35,2))
((Lisa,grade-2),(122,2))
((Mathew,grade-3),(45,1))
((Andrew,grade-2),(77,1))
((Andrew,grade-1),(131,3))
((Lisa,grade-3),(86,1))
((John,grade-1),(116,3))
((John,grade-2),(74,1))
((Mark,grade-1),(168,2))
((Andrew,grade-3),(70,2))
((Mathew,grade-2),(197,3))
```

Now we are calculating average by summing the marks and dividing by its count for each key. Below screenshot shows the final result,

```
scala> val StudAvg = RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}
StudAvg: org.apache.spark.rdd.RDD[(String, String), Double] = MapPartitionsRDD[188] at mapValues at <console>:30
scala> StudAvg.foreach(println)
((Lisa,grade-1),24.0)
((Mark,grade-2),17.5)
((Lisa,grade-2),61.0)
((Mathew,grade-3),45.0)
((Andrew,grade-2),77.0)
((Andrew,grade-1),43.666666666666664)
((Lisa,grade-3),86.0)
((John,grade-1),38.666666666666664)
((John,grade-2),74.0)
((Mark,grade-1),84.0)
((Andrew,grade-3),35.0)
((Mathew,grade-2),65.666666666666667)
```

Task 2.3 - What is the average score of students in each subject across all grades?

Codes,

```
val baseRDD =
```

```
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")  
(0),x.split(",")(1)),x.split(",")(3).toInt))
```

```
val RDDmap = baseRDD.mapValues(x=>(x,1))
```

```
val RDDreduce = RDDmap.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))
```

```
val SubAvg = RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}
```

We are first creating baseRDD to read the text file and we are extracting name and subject as key and marks as value,

```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")(0),x.split(",")(1)),x.split(",")(3).toInt))
18/01/03 02:54:11 WARN SizeEstimator: Failed to check whether UseCompressedOops is set; assuming yes
baseRDD: org.apache.spark.rdd.RDD[(String, String), Int] = MapPartitionsRDD[2] at map at <console>:24
```




ACADGILD

```
scala> baseRDD.foreach(println)
((Mathew,science),45)
((Mathew,history),55)
((Mark,maths),23)
((Mark,science),76)
((John,history),14)
((John,maths),74)
((Lisa,science),24)
((Lisa,history),86)
((Andrew,maths),34)
((Andrew,science),26)
((Andrew,history),74)
((Mathew,science),55)
((Mathew,history),87)
((Mark,maths),92)
((Mark,science),12)
((John,history),67)
((John,maths),35)
((Lisa,science),24)
((Lisa,history),98)
((Andrew,maths),23)
((Andrew,science),44)
((Andrew,history),77)
```

Now using **mapValues** we are mapping each value with 1-

```
scala> val RDDmap = baseRDD.mapValues(x=>(x,1))
RDDmap: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = MapPartitionsRDD[3] at mapValues at <console>:26
```

```
scala> RDDmap.foreach(println)
((Mathew,science),(45,1))
((Mathew,history),(55,1))
((Mark,maths),(23,1))
((Mark,science),(76,1))
((John,history),(14,1))
((John,maths),(74,1))
((Lisa,science),(24,1))
((Lisa,history),(86,1))
((Andrew,maths),(34,1))
((Andrew,science),(26,1))
((Andrew,history),(74,1))
((Mathew,science),(55,1))
((Mathew,history),(87,1))
((Mark,maths),(92,1))
((Mark,science),(12,1))
((John,history),(67,1))
((John,maths),(35,1))
((Lisa,science),(24,1))
((Lisa,history),(98,1))
((Andrew,maths),(23,1))
((Andrew,science),(44,1))
((Andrew,history),(77,1))
```



BIG DATA DEVELOPER ACADGILD Now we are adding the marks and number of occurrences for each key using **reduceByKey** and calculating average by dividing the sum of marks and count of occurrences for each key.

```
scala> val RDDreduce = RDDmap.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))
RDDreduce: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = ShuffledRDD[4] at reduceByKey at <console>:28

scala> RDDreduce.foreach(println)
((Lisa,history),(184,2))
((Mark,maths),(115,2))
((Andrew,science),(70,2))
((Mark,science),(88,2))
((Mathew,science),(100,2))
((Andrew,maths),(57,2))
((Mathew,history),(142,2))
((John,maths),(109,2))
((John,history),(81,2))
((Lisa,science),(48,2))
((Andrew,history),(151,2))

scala> val SubAvg = RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}
SubAvg: org.apache.spark.rdd.RDD[(String, String), Double] = MapPartitionsRDD[5] at mapValues at <console>:30

scala>

scala> SubAvg.foreach(println)
((Lisa,history),92.0)
((Mark,maths),57.5)
((Andrew,science),35.0)
((Mark,science),44.0)
((Mathew,science),50.0)
((Andrew,maths),28.5)
((Mathew,history),71.0)
((John,maths),54.5)
((John,history),40.5)
((Lisa,science),24.0)
((Andrew,history),75.5)
```

Task 2.4 - What is the average score of students in each subject per grade?

Codes,

val baseRDD =

```
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")
(1),x.split(",")(2)),x.split(",")(3).toInt))
```

val RDDmapvalue = baseRDD.mapValues(x=>(x,1))

val RDDreduce = RDDmapvalue.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))

val Avg_Grade =

```
RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}.foreach(println)
```

In first step we are creating paired RDD named as **baseRDD** to read the file and extracting subject and grade as key and marks as value



```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")(1),x.split(",")(2)),x.split(",")(3).toInt))
baseRDD: org.apache.spark.rdd.RDD[(String, String), (Int)] = MapPartitionsRDD[8] at map at <console>:24

scala> baseRDD.foreach(println)
((science,grade-3),45)
((history,grade-2),55)
((maths,grade-2),23)
((science,grade-1),76)
((history,grade-1),14)
((maths,grade-2),74)
((science,grade-1),24)
((history,grade-3),86)
((maths,grade-1),34)
((science,grade-3),26)
((history,grade-1),74)
((science,grade-2),55)
((history,grade-2),87)
((maths,grade-1),92)
((science,grade-2),12)
((history,grade-1),67)
((maths,grade-1),35)
((science,grade-2),24)
((history,grade-2),98)
((maths,grade-1),23)
((science,grade-3),44)
((history,grade-2),77)
```

Then we are mapping the values of **baseRDD** with numerical value 1 using function **mapValues**

```
scala> val RDDmapvalue = baseRDD.mapValues(x=>(x,1))
RDDmapvalue: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = MapPartitionsRDD[9] at mapValues at <console>:26

scala> RDDmapvalue.foreach(println)
((science,grade-3),(45,1))
((history,grade-2),(55,1))
((maths,grade-2),(23,1))
((science,grade-1),(76,1))
((history,grade-1),(14,1))
((maths,grade-2),(74,1))
((science,grade-1),(24,1))
((history,grade-3),(86,1))
((maths,grade-1),(34,1))
((science,grade-3),(26,1))
((history,grade-1),(74,1))
((science,grade-2),(55,1))
((history,grade-2),(87,1))
((maths,grade-1),(92,1))
((science,grade-2),(12,1))
((history,grade-1),(67,1))
((maths,grade-1),(35,1))
((science,grade-2),(24,1))
((history,grade-2),(98,1))
((maths,grade-1),(23,1))
((science,grade-3),(44,1))
((history,grade-2),(77,1))
```

We are adding marks and number of occurrences for each key using **reduceByKey**

```
scala> val RDDreduce = RDDmapvalue.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))
RDDreduce: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = ShuffledRDD[10] at reduceByKey at <console>:28

scala> RDDreduce.foreach(println)
((history,grade-2),(317,4))
((history,grade-3),(86,1))
((maths,grade-1),(184,4))
((science,grade-3),(115,3))
((science,grade-1),(100,2))
((science,grade-2),(91,3))
((history,grade-1),(155,3))
((maths,grade-2),(97,2))
```

We are calculating average by dividing the sum of marks with number of occurrences and the requested output is shown below,



```
scala> val Avg_Grade = RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}.foreach(println)
((history,grade-3),79.25)
((history,grade-3),86.0)
((maths,grade-1),46.0)
((science,grade-3),38.333333333333336)
((science,grade-1),50.0)
((science,grade-2),30.333333333333332)
((history,grade-1),51.666666666666664)
((maths,grade-2),48.5)
Avg_Grade: Unit = ()
```

Task 2.5 - for all students in grade-2, how many have average score greater than 50?

Codes,

val baseRDD =

```
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")(0),x.split(",")(2)),x.split(",")(3).toInt))
```

val RDDmap = baseRDD.mapValues(x=>(x,1))

val RDDreduce = RDDmap.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))

val RDDavg = RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}

val RDDfiltermap = RDDavg.filter(x=>x._1._2 == "grade-2" && x._2>50).count()

val RDDfiltermap = RDDavg.filter(x=>x._1._2 == "grade-2" && x._2>50).foreach(println)

Creating a paired RDD named as baseRDD to read the file and extracting name and grade as key and marks as value,

```
scala> val baseRDD = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")(0),x.split(",")(2)),x.split(",")(3).toInt))
baseRDD: org.apache.spark.rdd.RDD[(String, String), Int] = MapPartitionsRDD[14] at map at <console>:24

scala> baseRDD.foreach(println)
((Mathew,grade-3),45)
((Mathew,grade-2),55)
((Mark,grade-2),23)
((Mark,grade-1),76)
((John,grade-1),14)
((John,grade-2),74)
((Lisa,grade-1),24)
((Lisa,grade-3),86)
((Andrew,grade-1),34)
((Andrew,grade-3),26)
((Andrew,grade-1),74)
((Mathew,grade-2),55)
((Mathew,grade-2),87)
((Mark,grade-1),92)
((Mark,grade-2),12)
((John,grade-1),67)
((John,grade-1),35)
((Lisa,grade-2),24)
((Lisa,grade-2),98)
((Andrew,grade-1),23)
((Andrew,grade-3),44)
((Andrew,grade-2),77)
```

Now we are mapping each value of **baseRDD** with 1 as shown below,



```
scala> val RDDmap = baseRDD.mapValues(x=>(x,1))
RDDmap: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = MapPartitionsRDD[15] at mapValues at <console>:26

scala> RDDmap.foreach(println)
((Mathew,grade-3),(45,1))
((Mathew,grade-2),(55,1))
((Mark,grade-2),(23,1))
((Mark,grade-1),(76,1))
((John,grade-1),(14,1))
((John,grade-2),(74,1))
((Lisa,grade-1),(24,1))
((Lisa,grade-3),(86,1))
((Andrew,grade-1),(34,1))
((Andrew,grade-3),(26,1))
((Andrew,grade-1),(74,1))
((Mathew,grade-2),(55,1))
((Mathew,grade-2),(87,1))
((Mark,grade-1),(92,1))
((Mark,grade-2),(12,1))
((John,grade-1),(67,1))
((John,grade-1),(35,1))
((Lisa,grade-2),(24,1))
((Lisa,grade-2),(98,1))
((Andrew,grade-1),(23,1))
((Andrew,grade-3),(44,1))
((Andrew,grade-2),(77,1))
```

We are adding the marks of subject and number of occurrences per key using **reduceByKey** function-

```
scala> val RDDreduce = RDDmap.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))
RDDreduce: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = ShuffledRDD[16] at reduceByKey at <console>:28

scala> RDDreduce.foreach(println)
((Lisa,grade-1),(24,1))
((Mark,grade-2),(35,2))
((Lisa,grade-2),(122,2))
((Mathew,grade-3),(45,1))
((Andrew,grade-2),(77,1))
((Andrew,grade-1),(131,3))
((Lisa,grade-3),(86,1))
((John,grade-1),(116,3))
((John,grade-2),(74,1))
((Mark,grade-1),(168,2))
((Andrew,grade-3),(70,2))
((Mathew,grade-2),(197,3))
```

Here we are calculating the average of each student,

```
scala> val RDDavg = RDDreduce.mapValues{case(sum,count)=>(1.0*sum)/count}
RDDavg: org.apache.spark.rdd.RDD[(String, String), (Double)] = MapPartitionsRDD[17] at mapValues at <console>:30

scala>
scala> RDDavg.foreach(println)
((Lisa,grade-1),24.0)
((Mark,grade-2),17.5)
((Lisa,grade-2),61.0)
((Mathew,grade-3),45.0)
((Andrew,grade-2),77.0)
((Andrew,grade-1),43.666666666666664)
((Lisa,grade-3),86.0)
((John,grade-1),38.666666666666664)
((John,grade-2),74.0)
((Mark,grade-1),84.0)
((Andrew,grade-3),35.0)
((Mathew,grade-2),65.66666666666667)
```

Now in below step we are filtering the above result with student belonging to **grade-2** and having marks **greater than 50**, the number of the counts and the data is shown in the below.

```
scala> val RDDfiltermap = RDDavg.filter(x=>x._1._2 == "grade-2" && x._2>50).count()
RDDfiltermap: Long = 4
```




```
scala> val RDDfiltermap = RDDavg.filter(x=>x._1._2 == "grade-2" && x._2>50).foreach(println)
((Lisa,grade-2),65.0)
((Andrew,grade-2),77.0)
((John,grade-2),74.0)
((Mathew,grade-2),65.66666666666667)
RDDfiltermap: Unit = ()
```

Spark Operations – Problem Statement 3

Are there any students in the college that satisfy the below criteria:

1. Average score per student_name across all grades is same as average score per student_name per grade

Solution- To find the solution of above problem we will first calculate average of each student across all grades i.e. irrespective of grade. Below is the code used to find the same

val baseRDD1 =

sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>(x.split(",")(0),x.split(",")(3).toInt))

val studAvg = baseRDD1.mapValues(x=>(x,1)).foreach(println)

scala> val studReduce = studAvg.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))

scala> val Avg_Stud = studReduce.mapValues{case (sum,count) => (1.0 * sum)/count}

We created a paired RDD named as baseRDD1 by extracting only name and marks,

```
scala> val baseRDD1 = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>(x.split(",")(0),x.split(",")(3).toInt))
baseRDD1: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[133] at map at <console>:24

scala> baseRDD1.foreach(println)
(Mathew,45)
(Mathew,55)
(Mark,23)
(Mark,76)
(John,14)
(John,74)
(Lisa,24)
(Lisa,86)
(Andrew,34)
(Andrew,26)
(Andrew,74)
(Mathew,55)
(Mathew,87)
(Mark,92)
(Mark,12)
(John,67)
(John,35)
(Lisa,24)
(Lisa,98)
(Andrew,23)
(Andrew,44)
(Andrew,77)
```

Then we are mapping each value of above RDD with 1,



```
scala> val studAvg = baseRDD1.mapValues(x=>(x,1)).foreach(println)
(Mathew,(45,1))
(Mathew,(55,1))
(Mark,(23,1))
(Mark,(76,1))
(John,(14,1))
(John,(74,1))
(Lisa,(24,1))
(Lisa,(86,1))
(Andrew,(34,1))
(Andrew,(26,1))
(Andrew,(74,1))
(Mathew,(55,1))
(Mathew,(87,1))
(Mark,(92,1))
(Mark,(12,1))
(John,(67,1))
(John,(35,1))
(Lisa,(24,1))
(Lisa,(98,1))
(Andrew,(23,1))
(Andrew,(44,1))
(Andrew,(77,1))
```

Then we are adding the marks and number of occurrences for each student using **reduceByKey** as shown below,

```
scala> val studReduce = studAvg.reduceByKey((x,y)=> (x._1+y._1,x._2+y._2))
studReduce: org.apache.spark.rdd.RDD[(String, (Int, Int))] = ShuffledRDD[136] at reduceByKey at <console>:28

scala> studReduce.foreach(println)
(Mark,(203,4))
(Andrew,(278,6))
(Mathew,(242,4))
(John,(190,4))
(Lisa,(232,4))
```

In below step we are calculating the average of each student,

```
scala> val Avg_Stud = studReduce.mapValues{case (sum,count) => (1.0 * sum)/count}
Avg_Stud: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[137] at mapValues at <console>:30

scala> Avg_Stud.foreach(println)
(Mark,50.75)
(Andrew,46.333333333333336)
(Mathew,60.5)
(John,47.5)
(Lisa,58.0)
```

Now the second step of this problem is to find the average of each student per grade. We have used below code to find,

val baseRDD2 =

```
sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")(0),x.split(",")
(2)),x.split(",")(3).toInt)).foreach(println)
```

val grade = baseRDD2.mapValues(x=>(x,1))

val gradeReduce = grade.reduceByKey((x,y)=> (x._1+y._1,x._2+y._2))



```
val gradeAvg = gradeReduce.mapValues{case(sum,count) =>
(1.0*sum)/count}
```

So first we are creating another paired RDD named as baseRDD2 by extracting name and grade as key and marks as value from the input file,

```
scala> val baseRDD2 = sc.textFile("/home/acadgild/hadoop/17.2_Dataset.txt").map(x=>((x.split(",")(0),x.split(",")(2)),x.split(",")(3).toInt)).f
oreach(println)
((Mathew,grade-3),45)
((Mathew,grade-2),55)
((Mark,grade-2),23)
((Mark,grade-1),76)
((John,grade-1),14)
((John,grade-2),74)
((Lisa,grade-1),24)
((Lisa,grade-3),86)
((Andrew,grade-1),34)
((Andrew,grade-3),26)
((Andrew,grade-1),74)
((Mathew,grade-2),55)
((Mathew,grade-2),87)
((Mark,grade-1),92)
((Mark,grade-2),12)
((John,grade-1),67)
((John,grade-1),35)
((Lisa,grade-2),24)
((Lisa,grade-2),98)
((Andrew,grade-1),23)
((Andrew,grade-3),44)
((Andrew,grade-2),77)
baseRDD2: Unit = ()
```

Then we are mapping each value of baseRDD2 with 1 using **mapValues** function

```
scala> val grade = baseRDD2.mapValues(x=>(x,1))
grade: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = MapPartitionsRDD[153] at mapValues at <console>:26

scala> grade.foreach(println)
((Mathew,grade-3),(45,1))
((Mathew,grade-2),(55,1))
((Mark,grade-2),(23,1))
((Mark,grade-1),(76,1))
((John,grade-1),(14,1))
((John,grade-2),(74,1))
((Lisa,grade-1),(24,1))
((Lisa,grade-3),(86,1))
((Andrew,grade-1),(34,1))
((Andrew,grade-3),(26,1))
((Andrew,grade-1),(74,1))
((Mathew,grade-2),(55,1))
((Mathew,grade-2),(87,1))
((Mark,grade-1),(92,1))
((Mark,grade-2),(12,1))
((John,grade-1),(67,1))
((John,grade-1),(35,1))
((Lisa,grade-2),(24,1))
((Lisa,grade-2),(98,1))
((Andrew,grade-1),(23,1))
((Andrew,grade-3),(44,1))
((Andrew,grade-2),(77,1))
```

Then we are adding the marks and number of occurrences of 1 for each key using **reduceByKey()** function

```
scala> val gradeReduce = grade.reduceByKey((x,y)=>(x._1+y._1,x._2+y._2))
gradeReduce: org.apache.spark.rdd.RDD[(String, String), (Int, Int)] = ShuffledRDD[154] at reduceByKey at <console>:28

scala> gradeReduce.foreach(println)
((Lisa,grade-1),(24,1))
((Mark,grade-2),(35,2))
((Lisa,grade-2),(122,2))
((Mathew,grade-3),(45,1))
((Andrew,grade-2),(77,1))
((Andrew,grade-1),(131,3))
((Lisa,grade-3),(86,1))
((John,grade-1),(116,3))
((John,grade-2),(74,1))
((Mark,grade-1),(168,2))
((Andrew,grade-3),(70,2))
((Mathew,grade-2),(197,3))
```

In below step we are calculating average of each key by dividing the sum of marks with the count



```
scala> val gradeAvg = gradeReduce.mapValues{case(sum,count) => (1.0*sum)/count}
gradeAvg: org.apache.spark.rdd.RDD[(String, String), Double] = MapPartitionsRDD[155] at mapValues at <console>:30

scala> gradeAvg.foreach(println)
((Lisa,grade-1),24.0)
((Mark,grade-2),17.5)
((Lisa,grade-2),61.0)
((Mathew,grade-3),45.0)
((Andrew,grade-2),77.0)
((Andrew,grade-1),43.666666666666664)
((Lisa,grade-3),86.0)
((John,grade-1),38.666666666666664)
((John,grade-2),74.0)
((Mark,grade-1),84.0)
((Andrew,grade-3),35.0)
((Mathew,grade-2),65.666666666666667)
```

Now to proceed further we are extracting name and marks from above RDD,

In below step we are using intersection function between flatgradeAvg and flatnameAvg rdd's to find whether any common student is there.

```
scala> val flatgradeAvg = gradeAvg.map(x=> x._1._1 + "," + x._2.toDouble)
```

```
scala> val commanval = flatgradeAvg.intersection(flatAvg_Stud)
```

```
scala> val flatgradeAvg = gradeAvg.map(x=> x._1._1 + "," + x._2.toDouble)
flatgradeAvg: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[156] at map at <console>:32

scala> flatgradeAvg.foreach(println)
Lisa,24.0
Mark,17.5
Lisa,61.0
Mathew,45.0
Andrew,77.0
Andrew,43.666666666666664
Lisa,86.0
John,38.666666666666664
John,74.0
Mark,84.0
Andrew,35.0
Mathew,65.666666666666667

scala>

scala> val flatAvg_Stud = Avg_Stud.map(x=>x._1 + "," + x._2)
flatAvg_Stud: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[157] at map at <console>:32

scala> flatAvg_Stud.foreach(println)
Mark,50.75
Andrew,46.333333333333336
Mathew,60.5
John,47.5
Lisa,58.0
```

So the command **comman.foreach(println)** shows that no common students are there having average score per student_name across all grades is same as average score per **student_name** per grade-

```
scala> val commanval = flatgradeAvg.intersection(flatAvg_Stud)
commanval: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[163] at intersection at <console>:44

scala> commanval.foreach(println)

scala>
```

